

ROBOTIC BELIEF REVISION UNDER INTERVENTION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study interactive world models. The motivating bet is: Revise robot beliefs only when interventions reveal a violated physical assumption.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For interactive world models, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b; 2024; 2025; 2026c; 2022). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.562	0.438
observed_only_wam	hidden_mode_shift	0.152	0.848
augmentation_as_data_wam	nominal	0.608	0.392
augmentation_as_data_wam	hidden_mode_shift	0.282	0.718
scalar_uncertainty_wam	nominal	0.660	0.340
scalar_uncertainty_wam	hidden_mode_shift	0.355	0.645
proposed_branch_mechanism	nominal	0.782	0.218
proposed_branch_mechanism	hidden_mode_shift	0.545	0.455

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Robotic Belief Revision Under Intervention argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Developing real-world evidence from real-world data: Transforming raw data into analytical datasets, 2022. <https://doi.org/10.1002/lrh2.10293>.
- Generative world explorer, 2024. <http://arxiv.org/abs/2411.11844v3>.
- Development of compositionality through interactive learning of language and action of robots, 2025. <https://doi.org/10.1126/scirobotics.adp0751>.
- Pi-vla: A symmetry-aware predictive and interactive vision-language-action framework for robust robotic manipulation, 2026a. <https://doi.org/10.20944/preprints202601.0682.v1>.
- Pi-vla: Adaptive symmetry-aware decision-making for long-horizon vision-language-action manipulation, 2026b. <https://doi.org/10.3390/sym18030394>.
- Acwm-phys: Investigating generalized physical interaction in action-conditioned video world models, 2026c. <http://arxiv.org/abs/2605.08567v2>.